



Enabling Cyber-Secure Sovereign AI

Nations will increasingly rely on AI.
AI needs infrastructure.
Nations need their infrastructure to be Sovereign.
Sovereign infrastructure needs to be secured from the electron to the token.
We are building a cyber-secured, sovereign AI cloud. That is ASGARD. That is FrontierOne.

Securing AI Sovereignty is now every nation's priority. Sovereign AI compute is projected to grow rapidly in 2026, with investments expected to reach over \$100 billion by the end of 2026 and \$600 billion by 2030. The US Cloud Act has further accelerated the repatriation of data from hyper scalers into hyperlocal, secure AI factories for use in sensitive corporate and government workloads.
This surge is additionally fueled by countries and organizations prioritizing data control, security, regulatory compliance (such as the EU AI Act), and the desire for independence from US-centric AI providers.
The blocker to Sovereign AI is trust. With the release of Mythos, and additional cyber-focused models emerging from China, enterprises and governments cannot deploy AI at scale into national functions — energy, finance, health — if they cannot trust the full AI Stack: Datacenter, Compute, Cloud, Storage, Models.
This creates a market deadlock. Datacenters and bare-metal providers in global markets want enterprise and government compute contracts. Enterprises and governments want secure compute. Neither side can transact — because the trust layer in between does not exist.
FrontierOne has brought together a world-class team of cybersecurity experts, infrastructure operators, and go-to-market leaders to solve this problem — building ASGARD, the Cyber-Secure Sovereign Cloud that makes sovereign AI inference trusted for the first time.
We are already moving. Active partnerships with compute providers and datacenters are in place across 4 key markets — Norway (NATO, NonEU), Holland (EU), Canada and Vietnam. We are scaling faster than any competitor because ASGARD solves the same trust problem in every market simultaneously. One platform. Any jurisdiction. Assured security.

Providing Security and Trust to the Global Sovereign AI Market

Securing sovereign compute means protecting data, control, and operations so the workload stays under the right jurisdiction, the right operators, and the right cryptographic boundaries.

The strongest pattern is to combine local control, hardened infrastructure, compute and model governance, with real time cyber apparatus to ensure AI applications and models are monitored in real time and threats are detected and remediated.

This assurance layer for sovereign AI compute does not yet exist. FrontierOne is building it. Every dollar of sovereign AI infrastructure spend flows through compute that requires assurance — ASGARD captures the security premium on each GPU-hour processed through the secure stack.

SIZE THE 2030 SOVEREIGN AI SECURITY LAYER?			
SOVEREIGN AI MARKET BY 2030	Y1 PARTNER ARR PER DC (NORWAY)	SECURITY PREMIUM PER GPU-HOUR	TARGET GROSS MARGIN

\$600B McKinsey 2026	\$33.7M illustrative Y3 range \$400–700M per partner at full ramp	\$5.6–7 70% uplift over ~\$8/GPU-hr baseline	75%+ asset-light software model
--	---	--	---




\$100B Deployed in sovereign AI compute in 2026 alone · Deloitte	6.2GW Declared sovereign DC capacity outside US · Public filings
2/3 Of all AI compute is now inference — and growing · Deloitte 2026	35× Higher inference throughput/MW from NVIDIA LPX · NVIDIA GTC 2026

ASGARD: The Secure Sovereign Cloud

The AI infrastructure attack surface has expanded dramatically. Threat actors — state-sponsored and criminal — now have the ability to manipulate the core infrastructure of an enterprise and nation through dynamically changing AI applications. Software is no longer deterministic. Sovereign data plays a critical role in the output and functioning of AI-native applications. And the emergence of Mythos-class AI systems — autonomous models capable of chaining zero-days across entire operating systems — all expand the threat model exponentially.

The infrastructure that sovereign AI runs on is now itself the primary target. State-level adversaries can exploit the GPU cluster before the model ever generates a response.

ASGARD is FrontierOne's response. It is not a security tool bolted onto a cloud — it is the core operating system that assures every layer of the AI compute stack is secured.



The Cyber-Secure Compute Token (CSC Token) is FrontierOne's core technical output from ASGARD — our Secure Sovereign AI Cloud.

The CSC Token is what allows governments and enterprises to procure sovereign AI inference with the same confidence they apply to standard cloud services — mapped to the core security standards in each country.

02 · TEAM

Built by Operators at the Intersection of AI, Security, and Government

FrontierOne is the Avengers Assemble of cybersecurity and AI. We have brought together a world class team including DARPA-grade defence engineers, Cambridge AI security researchers, CISO veterans who have secured 30+ governments, HPC infrastructure leaders, and seasoned operators who have built and exited companies at the highest levels of tech and capital markets. Each person on this team is world-class in their domain.

Tee Ganbold

CEO



Cambridge Analytica Business Lead (age 23). Hired directly by Ethereum co-founder Joe Lubin to lead Business Development at ConsenSys at 25. CMO and Co-Founder at Clear AI.

Frank Meehan

CSO



Government/Sovereign Relations, Nebius, 2024–2025. Co-founder Equilibrium AI acquired by FiscalNote (NYSE IPO 2022). Board Director at Siri and Spotify (2009–2012). Led early investment into DeepMind. Investor in Anthropic, and Groq. Director of Engineering at Three UK.

Amit Gupta

CTO



Architect of FrontierOne's technical strategy. Scaled core engineering at Benchling (\$6B) and Lookout (\$2B). Deep expertise in secure distributed systems, AI infrastructure, and HPC. Drives ASGARD architecture, NVIDIA LPX integration, and the CSC Token engine.

Zach Powers

CISO



Built enterprise security programmes from the ground up at Benchling and One Medical. Previously VP Security at Salesforce, serving 30+ governments worldwide. Leads ASGARD's compliance posture: FedRAMP, CMMC, EUCS, IRAP, BSI C5, DEF-STAN.

Sameer
Chishty

CFO



Co-Founder of AsiaPak Investments (\$4B platform). Investor in OpenAI, Anthropic, and Groq. Former Partner at McKinsey (New York) and Morgan Stanley (Asia).

Jesse
Coffey

PRODUCT / ENGINEERING



Defence and aerospace platforms for DARPA, AFRL, and SPAWAR. Cyber challenge developer at DEFCON with the US Air Force Research Lab. Engineering leadership at Benchling (\$6B). Brings classified-grade engineering discipline to ASGARD's physical and firmware security layers.

Stacy
Gorelik

HPC / SOFTWARE LEAD



VP Engineering at Seqera (HPC bioinformatics). Head of Engineering at Flatiron Health. 20+ years delivering mission-critical software across financial services and health-tech. Leads ASGARD's compute orchestration and datacenter partner onboarding.

Colin
Rand

SECURITY / SOFTWARE LEAD



EVP R&D at SonicWall (Gartner SASE Magic Quadrant). Former VP Engineering at Banyan Security, leading zero-trust SSE and AI-powered defence against spearphishing. Prior VP Platform Engineering at Delphix and security engineering at Lookout (\$2B).

Dr.
David
Khachaturov

LEAD AI SECURITY



Cambridge PhD in adversarial machine learning. Harding Distinguished Scholar. Patented IP in computer vision security. Bye-Fellow at St Catharine's College, Cambridge. Leads ASGARD's AI threat modelling, model verification, and red-teaming.

Ricardo
Velez

CHIEF LEGAL OFFICER



Former Deputy General Counsel at Skype. General Counsel at Benchling (\$6B). Capital markets, M&A, and regulatory compliance across sovereign and enterprise technology environments.

Paul
Dufays

HEAD OF FINANCING



Cyber-Secure Cloud Business Model

The Sovereign AI market is being held back by a lack of trust and assurance. Enterprises are doing lots of POCs, but are afraid of rolling them out across the organisation because of lack of trust in the security around the POC models and applications. Datacenters and Compute Providers targeting the Sovereign AI market are not getting the long term enterprise contracts they built their financing around, because of this "Fear of Rollout". To solve this we:

- Partner with Datacenters and Compute Providers to provide the cyber-security assurance ASGARD on top on their GPUs and co-sell, deploy and manage ASGARD to enterprises and governments.
- Our partners share in the compute revenues (70% to FrontierOne, 30% to Partner) providing a 30% revenue uplift on their bare metal prices, together with long term enterprise contracts.
- Our partners have deep relationships in their markets with enterprises and governments, which provides the customer pipeline, and ASGARD is the catalyst to unlocking these customers.

FrontierOne/ASGARD contracts with the customer, deploys and manages the Sovereign AI solutions they need, which run on our partners infrastructure.

SECURITY PREMIUM OPPORTUNITY

\$5.6–7/GPU/hr at scale

Sovereign customers pay a verified security premium above bare metal GPU rates. At assured pipeline across 4 markets, the annual recurring revenue potential is transformative — recurring on every GPU-hour processed through ASGARD.

WHY FRONTIERONE SCALES FAST

Asset-light. No datacenter CapEx.

We assure partner infrastructure. We don't build it. Capital goes into software, accreditation, and business development — not concrete. This is why we can operate in 4 sovereign markets simultaneously while a closed-ecosystem competitor enters one per year.

TARGET GROSS MARGIN 75%+ Asset-light software model	PARTNER COMMITMENT FEE (PER DC) \$5M \$1M/MW cap · at signing	Y1 ARR PER PARTNER (NORWAY) \$33.7M Scaling to \$708M Y3
---	---	--

Targeting Enterprise and Government Secure Inference with our Datacenter Partners in 2026

Datacenters in Norway, Holland, Canada, and Vietnam are approaching FrontierOne because their enterprise and government customers require assured compute. We are their answer. This is an inbound-led go-to-market anchored in two complementary motions.

Datacenter Partner Assurance

We assure existing datacenter infrastructure, enabling partner facilities to access sovereign government and enterprise procurement they cannot reach alone. The partner pays a commitment fee at signing and shares inference revenue upside. We provide ASGARD deployment, accreditation support, and sales access to our sovereign customer pipeline.

ASGARD Secure Cloud

Government departments, defence ministries, critical infrastructure operators, and regulated enterprises in our 4 active markets procure directly through FrontierOne's ASGARD Secure Cloud — a fully managed, cyber-secure sovereign inference environment. These are multi-year, recurring contracts with high switching costs, anchored in the security assurances (FedRAMP, EUCS, IRAP, BSI C5, DEF-STAN) that no other provider has assembled simultaneously.

05 · COMPUTE PARTNER PIPELINE

Our Initial Compute Partners

PARTNER	MARKET	STATUS	HW AVAILABLE
Viking Digital	Norway	Final negotiations. Norwegian Government defence & health departments.	300MW
BW Velora	Norway & Canada	Final negotiations. Norway energy & shipping sector and Canadian sovereign AI infrastructure. BW Group-backed.	120MW
Digital Energy	Netherlands	Engaged. Shell, Adyen, Dutch enterprise pipeline.	30MW (planned)
GlobusAccess	Vietnam	Discovery stage. Government departments identified. 30MW datacenter target.	30MW (planned)

06 · BUSINESS MODEL

Secured Sovereign AI Inference

PRIMARY · PARTNER-OWNED INFRASTRUCTURE

70% Sov AI cloud revenues

Datacenter owns the bare metal inference chips. FrontierOne deploys ASGARD on top and captures a 70% security premium upside on every GPU-hour of sovereign compute processed — recurring, asset-light revenue with zero hardware risk.

SECONDARY · CO-OWNED INFRASTRUCTURE

Additional bare metal

FrontierOne participates in the bare metal revenue in cases where it facilitates the financing of compute chips with our datacenter partners.

TARGET GROSS MARGIN

75%+

Asset-light software model. No datacenter CapEx required to begin. Scales with megawatts / compute secured.


The model scales with the AI compute market buildout. Every new gigawatt of sovereign AI capacity deployed anywhere in our 4 markets is a potential assurance opportunity. The addressable market grows automatically with global AI infrastructure spend without having to finance the infrastructure CapEx.


The First Cyber-Secure Sovereign AI Cloud

The Sovereign AI compute market will explode in 2026, once cyber-security provides the assurance required. Agentic AI, real-time decision systems, and sovereign AI factories are all being built by countries and enterprises, yet the existing Sovereign AI Cloud solutions don't have the security AND flexibility required. The market is currently the following sectors:

- 1. Some companies are tightly integrated Secure Sovereign AI solutions, but they require the customer to use only their models. That isn't flexible enough.
- 2. While others have focused on building in country models, but don't have the security required.
- 3. While the NEO Clouds, like Nebius, focus on AI startups and don't have the on-premise, full stack security solutions required for enterprises.

It is our analysis that the company that will dominate the SovereignAI sector is the one focuses primarily on delivering the most secure cybersecurity Sovereign cloud, plus the ability to deploy on-premises for enterprise and government customers while not tied up in datacenter financing — **That is precisely where ASGARD operates.**





Full-stack cyber-security: hardware-rooted attestation and continuous threat detection across the entire compute stack — data center, compute, cloud, data, model — verified per transaction. Sovereign capability: ability to deploy in-country, on-prem, or air-gapped, fully outside the jurisdiction of foreign governance regimes including the US CLOUD Act.

08 · THE ROUND

\$15M Raise.
Use of Funds.

This round enables FrontierOne to complete the ASGARD core platform and close our initial datacenter partnerships (BW Velora, Digital Energy, Viking and Vietnam), and secure our first set of secure off-take customers using ASGARD compute. The objective is to deploy ASGARD in-country and produce the first live CSC Token on sovereign infrastructure.

\$8.0M · 53.3%	ASGARD security research, engineering, and product development. Forward deployment engineers — technical onboarding for committed and engaged partners. Digital Energy and BW Velora immediate.
\$3.0M · 20%	Business development and sales — government relations, sovereign contract BD across Norway, Holland, Canada, and Vietnam.
\$2.5M · 16.7%	Compliance and accreditation — EU Standards, CMMC L2, ISO 27001, UK Cyber Essentials+, IRAP, BSI C5. Global threat feeds, security tools, and operations. Compute costs
\$1.5M · 10%	Operations, G&A, Travel